

As AI Comes For
Your Data, Is Your
Recovery Ready?



Speaker



Jim Jones

**Sr. Product Infrastructure Architect,
Tl:tl Systems**

@k00laidIT

<https://koolaid.info>

Materials available at:

github.com/k00laidIT/Presentations/ElevateIT

Agenda

AI State of Play

- Capabilities driving adoption, corporate pressure, and data leverage

AI Putting Your Data At Risk


- Rapid deployment dangers, data leakage, and accelerated attacks

Building Cyber Resilience

- Governance, access controls, and immutable backup strategy

Questions held to the end

- Resources shared post-session



AI STATE OF PLAY

Capabilities | Corporate Pressure | Data as Force Multiplier

AI Is Rapidly Evolving

CODE DEVELOPMENT

90%+ of developers now use AI coding assistants. Untested, AI-generated code reaches production daily.

DOCS AND CONTENT

5-10x increase in speed with writing and creating documentation. Quality control and compliance accuracy are declining.

ASSISTANT- STYLE WORK

Customer service, research, and knowledge work are automated at scale. AI making decisions on corporate data.


Corporate Pressure to Deploy AI

- **Profitability Mandate**
Faster output, reduced headcount, margin pressure from leadership
- **Existing Contracts**
Employees required to use AI to achieve value from purchases
- **Long-Sought Efficiency**
The data already exists; AI just connects dots that humans couldn't, or couldn't do as quickly
- **Competitive Fear**
Peers' AI success creates urgency: deploy now, govern later (a dangerous pattern)



Your Data Is a Force Multiplier

- Historical data and today's data creation means exponential information growth
- Collaboration platforms (Slack, Teams, SharePoint) add noise, but also very useful information
- Your next competitive advantage is already in your data, waiting to be found
- AI extracts patterns humans miss:
AI is an opportunity and threat.



AI PUTTING YOUR DATA AT RISK

AI Security: What the Headlines Are Telling Us



AWS

THE RESILIENT CLOUD PLATFORM



McKinsey



Meta

II:II SYSTEMS

Rapid Deployment = Unintended Consequences

- **Policy Lag**
AI capabilities evolve faster than governance, compliance, and IT policy
- **Self-Healing Logic**
Models attempt to “clean up” after failures, triggering unintended data access and deletion
- **No Human in the Loop**
Agentic AI (e.g., OpenClaw-style agents) making production decisions autonomously, without oversight
- **Shadow Deployments**
Employees installing AI tools on personal devices; corporate data leaving the building silently

Alexey On Data

How I Dropped Our Production Database and Now Pay 10% More for AWS

A Terraform command executed by an AI agent wiped the production infrastructure behind the DataTalks.Club course platform. Here's how it happened and how the database was restored.



ALEXEY GRIGOREV

MAR 06, 2026

150

31

12

Share

I'm working on expanding the [AI Shipping Labs website](#) and wanted to migrate its current version from static GitHub Pages to AWS. And later, replace the original Next.js setup with a Django version.

My gradual plan was:

1. Move the current static site from GitHub Pages to AWS S3
2. Move DNS to AWS so the domain is fully managed there
3. Deploy the new Django version on a subdomain
4. When everything works, switch the main domain to Django

This way, everything would already be inside AWS, and the final switch would be seamless.

The migration strategy itself was reasonable, but the problems came from how I executed it.

Data Leakage Through Good Intentions

The Control Problem

- Improper AI integration to core data sources such as SaaS
- ACLs designed for humans, not AI agents with broad scope
- No data classification: AI accesses everything it can reach
- Compliance built on trust, not detection or enforcement

The Hacker News



🏠 Home

✉ Newsletter

🛒 Webinars

Microsoft Finds “Summarize with AI” Prompts Manipulating Chatbot Recommendations

👤 Ravie Lakshmanan 📅 Feb 17, 2026

Summarize with AI



Cybersecurity researchers disclosed they have detected a case of an information stealer infection successfully exfiltrating a victim's [OpenClaw](#) (formerly **Clawdbot** and **Moltbot**) configuration environment.

"This finding marks a significant milestone in the evolution of infostealer behavior: the transition from stealing [browser](#) credentials to harvesting the 'souls' and identities of personal AI [artificial intelligence] agents," Hudson Rock [said](#).

Alon Gal, CTO of Hudson Rock, told The Hacker News that the stealer was likely a variant of Vidar based on the infection details. Vidar is an [off-the-shelf information stealer](#) that's known to be active since late 2018.

Traditional Bad Actors: More Threats, Faster

24x7 AUTOMATION

AI = tireless script kiddies. Attacks run continuously while admins sleep. Vulnerability windows shrink from days to seconds.

NEW TARGETS

- Ransomware now targets AI models directly.
- Backup systems are the new attack surface.

RAPID ADAPTATION

- AI repeatedly transforms attacks to avoid detection making signature-based defenses obsolete. Detection always lags attack speed.

Over-reliance on Tooling

- AI enabled hotness
 - DSPM
 - DLP
 - AppSec
 - XDR
- Do you have the time to watch and manage all of these?
- Do you trust AI enabled dashboards to not hallucinate or miss things?

Paradigm Technica
Insights, Advice, and Solutions

Home Services Events About Us Contact Us Blog

1
Apr, 2026

The Claude Code Leak Isn't an Anthropic Problem. It's a Security Industry Problem.

The internet loves a stumble, and Anthropic handed the crowd a spectacular one on March 31, 2026. Pundits will queue up to dissect every embarrassing detail, competitors will quietly clone the feature roadmap they never paid for, and security professionals will perform the ritual of public tsk-tsking while privately feeling a warm current of schadenfreude. After all, this is the "safety-first" AI lab — the company that positions itself as the conscience of artificial intelligence — and it just leaked its own crown jewels to the public npm registry. The irony practically publishes itself — which, as it turns out, is exactly the problem.

But beneath the mockery lies a more important story, one that the cybersecurity industry would rather not confront.

What Actually Happened

When Anthropic shipped Claude Code version 2.1.88 to the public npm registry, it accidentally included a source map file — the kind of file build toolchains generate to help developers debug by connecting minified production code back to the original readable source. That map file contained a reference to an unobfuscated TypeScript archive hosted on Anthropic's own Cloudflare R2 storage bucket. Security researcher Chaofan Shou spotted it within hours, and the internet did what the internet does. His post on X amassed more than 28.8 million views, and the leaked codebase surpassed 84,000 stars and 82,000 forks on GitHub.

The zip archive contained nearly 1,900 TypeScript files and more than 512,000 lines of code — the full source of Claude Code, complete with internal libraries, slash commands, built-in tools, and unreleased feature flags. Those feature flags revealed a clear picture of how Anthropic is building toward longer autonomous tasks, deeper memory, and multi-agent collaboration — the precise engineering roadmap every competitor needed and nobody paid for. Anthropic confirmed the incident without ambiguity: "This was a release packaging issue caused by human error, not a security breach."

A Classic Insider Threat That No Tool Would Have Caught

<https://paradigmtechnica.com/2026/04/01/the-claude-code-leak-isnt-an-anthropic-problem-its-a-security-industry-problem/>



BUILDING CYBER RESILIENCE

THE RESILIENT CLOUD PLATFORM

 II:II SYSTEMS



Lock Down Your Data Protection Strategy

- **Audit**
Inventory all data sources, AI tool access, and ACLs. Do this today!
- **Define Retention**
Not all data needs to live forever. Archival strategy reduces your attack surface.
- **Day 0-2 Tools**
DR and backup solutions must be immediately accessible when an incident occurs.
- **Compliance by Design**
Governance must reflect real-world data flow, not wishful architecture.
- **Test, Test, Test**
Run incident response drills with AI threat scenarios. If you haven't tested recovery, you don't have recovery.

Cyber Resilience – Risk Controls Vulnerability Assessment

POD Rating	Cyber Resilience							
	Preventive Risk Controls (Pods)				Reactive Risk Controls (Pods)			
	Offensive Security	Cyber Hygiene	Threat Intelligence	Defensive Security	Monitor & Detect	Incident Response	Offline Backups	Incident Recovery
	Strong	Moderate	Moderate	Strong	Moderate	Very Weak	Very Weak	Very Weak
	Red & Purple Teaming	Train & Awareness	Dark Web Monitoring	Firewalls	IDS	Incident Prepared	Material Data	Clean Room
	Penetration Testing	Email Usage	Threat Profile	Intrusion Protect Sys	EDR	Response & Support	Offline Backups	Forensics
	Social Engineering	Mobile Devices	Strategic Intel	Identify Access Mgmt.	MDR	Detect & Contain	Immutability	Clean Data ID
	Vulnerability Scans	Remote Worksite	Technical Intel	Priv. Access Mgmt.	XDR	Forensics	Encryption	Recovery Execution
	Cyber TTX	Strong Passwords	Operational Intel	Segmentation	Cloud Monitoring	Restoration & Remed	Authentication	Plans & Proc.
	Risk Assessment	Multi-Factor Auth.	IOT Intel	Anti-Virus / Malware	SIEM	CM Support	Cyber Vault / AG	Lifecycle Mgmt
	TBD / NA	Device Currency	Geopolitical Intel	Patch Mgmt.	UAB	Cyber Insured Support	Phys. Secured	BC Plans
TBD / NA	Inti-virus SW	Indicator-based	System Hardening	SOAR	That Actor Eng.	Data Scanning	Recovery Testing	
TBD / NA	TBD / NA	TBD / NA	Encryption	TBD / NA	TBD / NA	TBD / NA	Clean HW / Systems	


Protect Production Data Through Access Controls



- **Principle of least privilege for every AI agent:** read-only, specific fields only
- **MCP and API gateway rules:** gate and log what AI can request
- **Real-time monitoring:** alert on unusual data access patterns immediately
- **Zero-trust for AI:** verify every data request, not just the connection

Immutable, Encrypted Backups for the AI Era

- **Immutability is non-negotiable:** backups cannot be modified or deleted after creation
- **Encryption at-rest and in-transit:** multi-layer protection as standard
- **Dynamic tagging and policy:** backups inherit data sensitivity and retention rules
- **Replicated across regions:** 11:11's 60+ global locations ensure geographic redundancy
- **DR systems isolated from production AI scope:** separate network, separate credentials





11:11 Cyber Vault for Cohesity | Veeam | Zerto

Storage Class	AWS S3 Standard Infrequent Access
Georedundancy	11-nines of durability
Read/restore	Unlimited
Data Locality Choice	25 global AWS regions
Immutability	Yes
Encryption	Yes
Minimum Retention	30 Days
Price	Low per TB contractable rate with no extra fees (API, egress, etc.)



Cyber Recovery Actions



Q&A

Materials available at:
github.com/k00laidIT/Presentations/ElevateIT

THANK YOU

111SYSTEMS.COM